

# Li, Li, Huo - Optimal In-Place Suffix Sorting

---

Patrick Indri

June 10, 2019

DSSC - Algorithmic Design Exam

- Problem setting;
- Naive solution;
- Preliminary notions;
- Suffix sorting for read-only integer alphabets;
- Additional results and conclusions;
- Auxiliary material.

*Suffix Arrays*: a space-saving alternative to suffix trees.

### Definition

Given a string  $T = T[0, \dots, n-1]$  where each  $T[i] \in \Sigma$  integer alphabet, the *suffix array*  $SA$  contains the indices of all suffixes of  $T$  which are sorted in lexicographical order.

*Suffix Arrays*: a space-saving alternative to suffix trees.

### Definition

Given a string  $T = T[0, \dots, n-1]$  where each  $T[i] \in \Sigma$  integer alphabet, the *suffix array*  $SA$  contains the indices of all suffixes of  $T$  which are sorted in lexicographical order.

### Example

$T = "1120"$ , the suffixes are  $\{1120, 220, 20, 0\}$ . Since  $\text{suf}(3) < \text{suf}(0) < \text{suf}(2) < \text{suf}(1)$ , then  $SA = [3021]$ .

## Problem

Construct  $SA$  for a given string  $T$ .

## Main Theorem

There is an in-place linear time algorithm for suffix sorting over integer alphabets, even if the input string  $T$  is read only and the size of the alphabet  $|\Sigma|$  is  $O(n)$ .

# Naive Solution

- Get all the suffixes and sort them using *Quicksort*, while retaining their original indices.  $O(n \log n)$  comparisons for sorting,  $O(n)$  to compare suffixes: worst case is  $O(n^2 \log n)$ .
- Build a suffix tree in  $O(n)$  and perform a depth-first traversal on it in  $O(n)$ .

## Notations:

- $\text{suf}(i)$  is said to be  $S$ -suffix if  $\text{suf}(i) < \text{suf}(i + 1)$ . Otherwise, it is  $L$ -suffix;
- $\text{suf}(i)$  is said to be  $LMS$ -suffix if  $\text{suf}(i)$  is  $S$ -suffix and  $\text{suf}(i - 1)$  is  $L$ -suffix;

## Note

Types ( $S$  or  $L$ , and  $LMS$ ) can be computed by a linear scan of  $T$ .

# Preliminary Notions

## Notations:

- $\text{suf}(i)$  is said to be  $S$ -suffix if  $\text{suf}(i) < \text{suf}(i + 1)$ . Otherwise, it is  $L$ -suffix;
- $\text{suf}(i)$  is said to be  $LMS$ -suffix if  $\text{suf}(i)$  is  $S$ -suffix and  $\text{suf}(i - 1)$  is  $L$ -suffix;

## Note

Types ( $S$  or  $L$ , and  $LMS$ ) can be computed by a linear scan of  $T$ .

**Example:**  $T = "31221120"$

Index	0	1	2	3	4	5	6	7
$T$	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*



# Suffix Sorting for Read-only Integer Alphabets

## Definitions and assumptions:

- $n_S$  ( $n_L$ ) denotes the number of  $S$ -suffixes ( $L$ -suffixes);
- $n_L \leq n_S$ ;
- $n_1$  denotes the number of  $LMS$ -suffixes;
- $SA[0, \dots, n - 1]$  will store the result;
- Bucket: set of suffixes with the same first character.

# Suffix Sorting for Read-only Integer Alphabets

## Definitions and assumptions:

- $n_S$  ( $n_L$ ) denotes the number of  $S$ -suffixes ( $L$ -suffixes);
- $n_L \leq n_S$ ;
- $n_1$  denotes the number of  $LMS$ -suffixes;
- $SA[0, \dots, n - 1]$  will store the result;
- Bucket: set of suffixes with the same first character.

## Algorithm:

1. Sort all  $LMS$ -characters of  $T$ ;
2. Induced sort all  $LMS$ -substrings from sorted  $LMS$ -characters;
3. Construct the reduced problem  $T_1$  from sorted  $LMS$ -substrings;
4. Sort the  $LMS$ -suffixes by recursively solving  $T_1$ ;
5. Induced sort all suffixes of  $T$  from the sorted  $LMS$ -suffixes.

# 1. Sort all *LMS*-characters of $T$

The *LMS*-characters can be sorted with *Counting Sort*, using  $SA[0, \dots, n/2]$  as the counting array and storing the result in  $SA[n - n_1, \dots, n - 1]$ . AUX

Index	0	1	2	3	4	5	6	7
$T$	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*

SA

						7	1	4
--	--	--	--	--	--	---	---	---

The sorting step takes  $O(n)$  time and uses  $O(1)$  workspace.

## 2. Induced sort all *LMS*-substrings from sorted *LMS*-characters

This step is analogous to step 5.

After this step, indices of the ordered *LMS*-substrings are stored in  $SA[n - n_1, \dots, n - 1]$ .

Index	0	1	2	3	4	5	6	7
T	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*

SA

						7	1	4
--	--	--	--	--	--	---	---	---



SA

						7	4	1
--	--	--	--	--	--	---	---	---

*LMS*-substrings are  $\{1121, 1120, 0\}$ .

### 3. Construct the reduced problem $T_1$

Using the lexicographical order of the *LMS*-substrings, build the reduced problem  $T_1$  and store it in  $T[0, \dots, n_1]$ .  $T_1$  can be obtained by a liner scan of *SA*, thus using  $O(n)$  time and  $O(1)$  workspace.

Index	0	1	2	3	4	5	6	7
T	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*

SA

					7	4	1
--	--	--	--	--	---	---	---



SA

2	1	0			7	4	1
---	---	---	--	--	---	---	---

*LMS*-substrings are  $\{1121, 1120, 0\}$ .

#### 4. Sort the *LMS*-suffixes by recursively solving $T_1$

$T_1$  can be solved iteratively in linear time<sup>1</sup> with no additional workspace. It is stored at the beginning of *SA*. The *LMS*-suffixes are sorted using linear scans and the solution of  $T_1$ .

Index	0	1	2	3	4	5	6	7
T	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*

SA	2	1	0					
	↓							
SA	2	1	0			1	4	7
	↓							
SA	7	4	1			1	4	7
	↓							
SA						7	4	1

*LMS*-suffixes are  $\{1221120, 1120, 0\}$ .

The complexity of this step is  $O(n)$  in time and  $O(1)$  in workspace.

---

<sup>1</sup> $\mathcal{T}(n) = \mathcal{T}(n/2) + n = n(1 + 1/2 + 1/4 + 1/8 + \dots + 1/\log_2 n) \in \Theta(n)$

## 5. Induced sort all suffixes of $T$ from the sorted $LMS$ -suffixes

It can be demonstrated that sorting the  $n_L$   $L$ -suffixes from the sorted  $LMS$ -suffixes is symmetrical as sorting the  $n_S$   $S$ -suffixes from the sorted  $L$ -suffixes. Suppose  $L$ -suffixes are sorted.

Index	0	1	2	3	4	5	6	7
T	3	1	2	2	1	1	2	0
Type	L	S	L	L	S	S	L	S
LMS		*			*			*

SA

						7	4	1
--	--	--	--	--	--	---	---	---



SA

6	3	2	0					
---	---	---	---	--	--	--	--	--

$L$ -suffixes are  $\{31221120, 221120, 21120, 20\}$ .

## 5. Induced sort all suffixes of $T$ from the sorted $LMS$ -suffixes

### Pointer Data Structure

Built in linear time, indicates the bucket tails of a  $S$ -suffix in constant time. Occupies at most  $c_P = cn / \log n$  words, placed at the end of  $SA$ . AUX

### Interior Counter Trick

Dynamically maintain the  $RF$ -pointers (rightmost free pointers) for each bucket. AUX



## 5. Induced sort all suffixes of $T$ from the sorted $LMS$ -suffixes

### Pointer Data Structure

Built in linear time, indicates the bucket tails of a  $S$ -suffix in constant time. Occupies at most  $c_P = cn / \log n$  words, placed at the end of  $SA$ . AUX

### Interior Counter Trick

Dynamically maintain the  $RF$ -pointers (rightmost free pointers) for each bucket. AUX

The ordering of the  $S$ -suffixes proceeds in two steps:

1. Construct a *pointer data structure*  $\mathcal{P}$  and, combining it with the *interior counter trick* induce the first  $n_S - c_P$   $S$ -suffixes;
2. Use *Binary Search* and the *Interior Counter Trick* on the last  $c_P$   $S$ -suffixes.

## 5. Induced Sort Algorithm

Suppose  $L$ -indices are already sorted in  $SA[0, \dots, n_L]$ .


```
def InducedSort(T, SA):  
    for i = nL downto 0:  
        j = SA[i] - 1 # Indicizes the bucket.  
        if T[j] is S-type:  
            *RF[j] = j  
            RF[j] = next_free_entry  
        endif  
    endfor  
enddef
```

For each query of  $RF$ , the tail of the bucket is provided in constant time by the pointer data structure, from that the interior counter trick indicates the  $RF$  entry.

## 5. Induced sort all suffixes of $T$ from the sorted $LMS$ -suffixes

The first  $n_S - c_P$   $S$ -suffixes can be ordered in linear time with no additional workspace. The remaining suffixes can be sorted without  $\mathcal{P}$ , using binary search to find the tails of the buckets: since  $c_P \log n = O(n)$ , time linearity is preserved.

A stable, in place linear time merging can be used to merge the sorted  $S$ - and  $L$ -suffixes.



## Additional results and conclusion

### **(Read-only) Integer Alphabets**

Considering all the steps, it follows that the algorithm takes  $O(n)$  time and  $O(1)$  workspace to compute the suffix array of a string  $T$  over integer alphabets  $\Sigma$ , where  $T$  is read-only and  $|\Sigma| = O(n)$ .

The result trivially holds for non-read-only integer alphabets.

## Additional results and conclusion

### **(Read-only) Integer Alphabets**

Considering all the steps, it follows that the algorithm takes  $O(n)$  time and  $O(1)$  workspace to compute the suffix array of a string  $T$  over integer alphabets  $\Sigma$ , where  $T$  is read-only and  $|\Sigma| = O(n)$ .

The result trivially holds for non-read-only integer alphabets.

### **Read Only General Alphabets**

For read-only general alphabets (i.e., only comparisons allowed on  $T$ ) there is an in-place  $O(n \log n)$  time algorithm for suffix sorting.

## Additional results and conclusion

### (Read-only) Integer Alphabets

Considering all the steps, it follows that the algorithm takes  $O(n)$  time and  $O(1)$  workspace to compute the suffix array of a string  $T$  over integer alphabets  $\Sigma$ , where  $T$  is read-only and  $|\Sigma| = O(n)$ .

The result trivially holds for non-read-only integer alphabets.

### Read Only General Alphabets

For read-only general alphabets (i.e., only comparisons allowed on  $T$ ) there is an in-place  $O(n \log n)$  time algorithm for suffix sorting.

### References

Li, Zhize, Jian Li, and Hongwei Huo. "Optimal in-place suffix sorting." *International Symposium on String Processing and Information Retrieval*. Springer, Cham, 2018.

## Auxiliary Material

---

## AUX: Sort all *LMS*-characters of $T$

Since  $|\Sigma| = O(n)$ , assume  $\exists d \in \mathbb{N}$  s.t.  $|\Sigma| \leq dn$ . Divide *LMS*-characters in  $2d$  partitions, where partition  $i$  contains elements in  $\left[ \frac{i|\Sigma|}{2d} + 1, \frac{(i+1)|\Sigma|}{2d} \right]$ . Since

$$\frac{|\Sigma|}{2d} \leq \frac{dn}{2d} = \frac{n}{2},$$

$SA[0, \dots, n/2]$  can be used as a counting array. 



## AUX: Interior counter trick - 1

Consider a bucket of size  $m$ , indexing it as  $SA_S\{0, \dots, m-1\}$ . Define the special symbols  $B_H$ ,  $B_T$ ,  $E$ ,  $R_1$  and  $R_2$ .  $\text{Index}(i)$  denotes the index of the  $i$ -th  $S$ -suffix of the bucket. The position of the tail, i.e.,  $m-1$ , is given by the pointer data structure.

```
def InteriorCounterTrick(SAs):  
    SAs[0]=BH, SAs[m-2]=E, SAs[m-1]=BT  
    # O(m) time.  
    if SAs[m-1]=BT and  
        (SAs[m-2]=E or SAs[m-SAs[m-2]-3]!=BH):  
        for i=1 upto m-3:  
            SAs[m-2-i]=Index(i)  
            SAs[m-2]++      # Acts as a counter.  
        endfor  
    endif
```

## AUX: Interior counter trick - 2

```
# O(m) time.  
if SAs[m-1]=BT and SAs[m-SAs[m-2]-3]=BH:  
    shift SAs[1,...,m-3] to SAs[2,...,m-2]  
    SAs[1]=Index(m-2)  
    SAs[m-1]=R2  
endif  
# O(m) time.  
if SAs[m-1]=R2;  
    shift SAs[1,...,m-2] to SAs[2,...,m-1]  
    SAs[1]=Index(m-1)  
    SAs[0]=R1  
endif
```

## AUX: Interior counter trick - 3

```
# O(m) time, need to scan from tail  
# backwards to find R1.  
else:  
    SAs[0] = Index(m)  
enddef
```

The function consists of four steps, each  $O(m)$  time, assuming that the tail of a bucket is known. It uses  $O(1)$  workspace and, for all the buckets, results in  $O(n)$  time. ↔

## AUX: Pointer data structure - 1

Assuming  $|\Sigma| \leq dn$ , divide the  $S$ -suffixes of  $T$  in  $4d$  parts, according to their first character. Let  $D_i$  denote the pointer data structure of the  $i$ -th part.  $D_0$  can be constructed as follows (analogously for the others). For brevity,  $b = |\Sigma|/4d$ .

```
def PointerDataStructure(T, SAs):  
    SAs[i]=1 forall i in [1,b]  
    for i=n-1 downto 0:  
        if T[i] is S-type and in [1,b]: SAs[T[i]]++  
    endfor
```

## AUX: Pointer data structure - 2

```
sum=-1
for i=1 upto b:
    sum+=SAs[i]
    SAs[i]=sum
endfor
enddef
```

For every  $S$ -suffix for which  $T[i] \in SA_S[i, |\Sigma|/4d]$ ,  
 $SA_S[T[i]] - T[i]$  indicates the tail of the bucket  $T[i]$ : the tail of a bucket can be obtained in constant time. 